

# Distributed Task-Based Training of Tree Models

Da Yan<sup>\*1</sup>, Md Mashiur Rahma Chowdhury<sup>\*2</sup>, Guimu Guo<sup>\*3</sup>, Jalal Kahlil<sup>\*4</sup>, Zhe Jiang<sup>†5</sup>, Sushil Prasad<sup>‡6</sup>

<sup>\*</sup>Department of Computer Science, The University of Alabama at Birmingham

{<sup>1</sup>yanda, <sup>2</sup>mashiur, <sup>3</sup>guimuguo, <sup>4</sup>jalalk}@uab.edu

<sup>†</sup>Department of Computer & Information Science & Engineering, University of Florida

<sup>5</sup>zhe.jiang@ufl.edu

<sup>‡</sup>Department of Computer Science, The University of Texas at San Antonio

<sup>6</sup>sushil.prasad@utsa.edu

**Abstract**—Decision trees and tree ensembles are popular supervised learning models on tabular data. Two recent research trends on tree models stand out: (1) bigger and deeper models with many trees, and (2) scalable distributed training frameworks. However, existing implementations on distributed systems are *IO-bound* leaving CPU cores underutilized. They also only find best node-splitting conditions *approximately* due to row-based data partitioning scheme. In this paper, we target the exact training of tree models by effectively utilizing the available CPU cores. The resulting system called *TreeServer* adopts a *column-based* data partitioning scheme to minimize communication, and a *node-centric task-based* engine to fully explore the CPU parallelism. Experiments show that *TreeServer* is up to 10× faster than models in Spark MLlib. We also showcase *TreeServer*’s high training throughput by using it to build big “deep forest” models.

## I. INTRODUCTION

Decision tree is a supervised model trained over a data table to predict the value of a target attribute  $Y$  for an entity based on the observed values of its other attributes  $A_1, A_2, \dots, A_m$ . For example, given a customer data table in Fig. 1(a), Fig. 1(b) shows a decision tree to predict whether someone may default on a credit card payment. To reach a decision ( $Y$ -value at a leaf), a list of questions on attributes like “Age” and “Income” are gone through starting from the root.

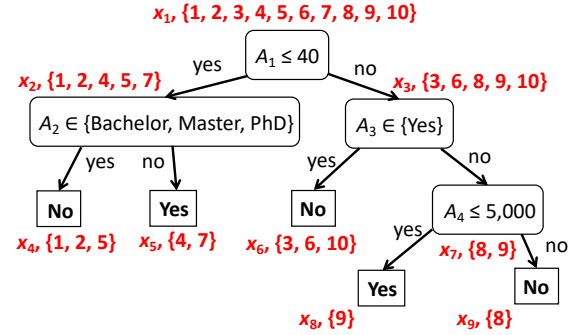
Recently, deeper ensemble models with many trees like deep forests [34] and mGBDTs [19] were shown to outperform deep neural networks in a number of tasks. Moreover, the distributed tree training algorithm, Google’s PLANET [27], has been implemented as a standard model in distributed data science libraries such as Mahout [26] and Spark MLlib [24].

However, since PLANET was originally proposed for MapReduce that keeps data on HDFS as row blocks, while nodes in tree models are split by conditions on attributes (i.e., columns), PLANET cannot support efficient computation of the best node-splitting conditions. Note that since the values of each individual column are stored in a distributed manner, excessive communication is necessary to find the best node-splitting conditions. To mitigate this issue, most existing implementations compute approximate node-splitting conditions, such as in Spark MLlib [24] and XGBoost [17].

Note that a Spark-based system such as MLlib cannot fundamentally address this problem, since RDDs are distributed

	$A_1$ : Age	$A_2$ : Highest Education	$A_3$ : Home Owner	$A_4$ : Income	$Y$ : Default
1	24	Bachelor	No	5,000	No
2	28	Master	Yes	7,500	No
3	44	Bachelor	Yes	5,500	No
4	32	Secondary	Yes	6,000	Yes
5	36	PhD	No	10,000	No
6	48	Bachelor	Yes	6,500	No
7	37	Secondary	No	3,000	Yes
8	42	Bachelor	No	6,000	No
9	54	Secondary	No	4,000	Yes
10	47	PhD	Yes	8,000	No

(a) Data Table  $D$



(b) Decision Tree  $\Delta_{x1}$

Fig. 1. Data Table and Tree Notations

in-memory row blocks. Even though Spark is switching to a “DataFrame” data organization to support column abstractions, the DataFrames are still row-based and actually immutable. New columns can be computed from DataFrames but they are still stored in a distributed manner, and a machine still cannot access an entire column without communication.

Another efficiency issue with PLANET is that, it adopts a top-down approach for node construction. Due to row-based data partitioning, upper-level node processing is IO-bound, making CPU utilization low initially during tree construction.

Recently, T-thinker [30] has been proposed as a parallel computing paradigm to overcome the problem of CPU underutilization in existing data-intensive Big Data systems, which were found to have a comparable or lower throughput than a single-threaded program [20], despite the use of aggregate IO throughput from many machines. T-thinker has seen success in graph mining problems [21], [22], [28], [29], [31]–[33].

T-thinker targets divide-and-conquer problems where a computing task over a big dataset can be recursively divided into independent tasks over smaller data subsets for parallel computation. Since the communication cost of collecting data for a task is linear to the data size, while computing workloads grow quickly with the data size, the latter cost can surpass the former as long as the data of a task is not too small. If a task is waiting for data, T-thinker suspends it to release a CPU core for use by other tasks, and the task will be timely resumed when its requested data are all received. This allows communication to be overlapped with computation to keep CPU cores busy with the actual computation.

Following the T-thinker paradigm, we propose *TreeServer*, a distributed system for training tree models where each task corresponds to the construction of a subtree rooted at a node  $x$ , denoted by  $\Delta_x$ ; this task can be divided into subtasks that construct subtrees rooted at  $x$ 's child nodes. If the set of training data that fall into  $x$  is small enough for processing by a machine, a task pulls these data to build the entire  $\Delta_x$  locally without additional communication to keep a core busy.

We identify three challenges in realizing the above idea: (1) how to schedule such CPU-bound tasks early to utilize the idle CPU cores in contrast to PLANET's top-down level-by-level node processing which delays effective CPU utilization? (2) how to compute the best attribute splitting conditions at each node exactly without excessive communication as experienced by current systems? (3) since we adopt a master-workers architecture where the master manages tree construction tasks, how to effectively direct the data-pulling requests from workers to different other workers without master involvement, so that data communication would not overwhelm any machine?

*TreeServer* provides elegant solutions to all the above 3 challenges. The main contributions of this paper are as follows:

- *TreeServer* partitions tabular data among machines by columns, which allows each attribute to be checked for node splitting without excessive communication.
- *TreeServer* schedules node-centric tasks in a hybrid scheme combining breadth-first with depth-first traversal, to schedule CPU-bound subtree construction tasks early.
- *TreeServer* keeps the data rows associated with different tree nodes in different machines to avoid a single point of communication-bottleneck for transmitting the subset of records requested by child-node tasks.
- Sound algorithms are developed to realize the above designs, plus a cost model for load balancing.

Besides the above desirable designs, *TreeServer* supports all types of attributes and handles missing data and attribute values unseen during training. *TreeServer* is fully compatible with the Hadoop ecosystem and loads data in parallel from Hadoop Distributed File System (HDFS). *TreeServer* is an ideal building block for training larger tree ensembles such as deep forests [34] in a Hadoop analytics workflow. Extensive experiments show that *TreeServer* is up to 10 $\times$  faster than Spark MLlib, and constructs deep forests efficiently.

The rest of this paper is organized as follows. Section II reviews our notations and the related work. Section III overviews

the design and features of *TreeServer*. Section IV describes the computation workflow of *TreeServer* threads. Following that, Section V presents our approach to release master from the duty of relaying row indices of nodes to workers to avoid delaying the transmission of task control messages, and Section VI introduces our approach to assign task workloads to worker machines. Finally, Section VII provides a case study of using *TreeServer* to construct deep forests, Section VIII reports our experiments, and Section IX concludes this paper.

To ensure reproducibility, we have released our code for *TreeServer* [11] and deep forest on top [13]. We also provide demo videos for *TreeServer* (short [10] and full [9] versions) and deep forest [12] with detailed steps to repeat experiments.

## II. PRELIMINARIES

**Notations and Decision Tree Review.** We now introduce our notations, and a complete list of notations used in this paper can be found in Appendix A for quick reference.

We consider a data table  $D$  with  $m$  attributes  $\mathcal{A} = \{A_1, A_2, \dots, A_m\}$  where one of them is the attribute  $Y$  to predict. In Fig. 1(a), we have  $m = 5$  and  $Y = A_5$ . We also denote the set of rows that fall into node  $x$  by  $D_x$ , and denote their IDs by  $I_x$ . In Fig. 1(b),  $I_{x_2} = \{1, 2, 4, 5, 7\}$ . Finally, we denote the subtree rooted at node  $x$  by  $\Delta_x$ .

During training, at each node  $x$ , we want to find a **split-condition** to partition the rows of  $D_x$  into child nodes such that an impurity score is reduced the most. We only consider binary node splitting since any multi-way splitting can be represented as a series of binary splitting. Each node  $x$  splits  $D_x$  based on an attribute  $A_i$  called the **split-attribute**.

There are 2 cases for the split-condition: (1) if  $A_i$  is ordinal, then the split-condition is " $A_i \leq v$ ", e.g., node  $x_1$  has  $A_1 \leq 40$  in Fig. 1(b); (2) if  $A_i$  is categorical, and its possible values constitute a set  $S_i$ , then the split-condition is " $A_i \in S_\ell$ " where  $S_\ell \subset S_i$ . If a row in  $D_x$  has an attribute value of  $A_i$  that falls in  $S_\ell$ , it goes to the left child node (right child otherwise). In Fig. 1, Attribute  $A_2$  has 5 possible values  $S_i = \{\text{Primary, Secondary, Bachelor, Master, PhD}\}$ , and  $S_\ell = \{\text{Bachelor, Master, PhD}\}$ . Note that if  $A_i$  is categorical, there are exponentially many possible split-conditions.

The split-attribute  $A_i$  and its split-condition (i.e.,  $v$  or  $S_\ell$ ) are greedily selected from a set of candidate attributes  $\mathcal{C} \subseteq \mathcal{A}$  to minimize impurity after shattering  $D_x$  into the two child nodes. The quality of node splitting is evaluated quantitatively using an impurity function such as the entropy of  $Y$ -labels for classification and variance of  $Y$ -values for regression. At each node  $x$ , one may randomly sample a subset  $\mathcal{C} \subseteq \mathcal{A}$  of attributes, and then greedily select the best attribute  $A_i \in \mathcal{C}$  and its split-condition. A special case is *completely random decision tree* [18] where to split each node, only one feature is randomly sampled (i.e.,  $|\mathcal{C}| = 1$ ). In other cases,  $\mathcal{C}$  may be sampled beforehand and then used throughout the building of a tree, as is in the ensemble model *random forest* where each tree is trained on a randomly sampled attribute subset.

At each tree node  $x$ , we select the best split-condition for each individual attribute (column)  $A_i$  independently (hence

can run in parallel), often just need one pass over the column values of  $A_i$  and  $Y$  for rows in  $D_x$  using well-known methods for decision tree training. The detailed algorithm can be found in Appendix B. The best conditions of different attributes  $A_i \in \mathcal{C}$  are then compared to select  $x$ 's overall best split-condition.

A node  $x$  becomes a leaf for reasons such as (1) all rows in  $D_x$  have the same  $Y$ -label and thus the prediction at the node is unique; (2)  $|D_x| \leq \tau_{leaf}$  where  $\tau_{leaf}$  is a user-defined threshold; (3) the **depth** of node  $x$ , denoted by  $d(x)$ , reaches the maximum allowed threshold  $d_{max}$ . Conditions (2) and (3) are to avoid overfitting the training data, and in classification, the label is predicted as the label that most rows in  $D_x$  have, or simply a PMF (probability mass function) vector over all labels. For example, in Fig. 1, if we stop at node  $x_7$  as a leaf, since  $D_{x_7}$  contains Row 8 with label “No” and 9 with label “Yes”,  $x_7$  will output a PMF vector {“Yes”: 50%, “No”: 50%} or a predicted label “Yes” (or “No” as we have a tie). For regression where  $Y$  is numerical, only conditions (2) and (3) applies, and the predicted value is the average  $Y$ -value of  $D_x$ .

**Related Systems.** Since MapReduce and Spark’s RDD partition input data among machines by rows, to find the best value  $v$  for a split-condition “ $A_i \leq v$ ” for a node  $x$ , every machine needs to check its own portion of those rows that belong to  $D_x$ . Since there can be up to  $n = |D_x|$  different  $A_i$ -values for rows in  $D_x$ , we need to consider up to  $(n - 1)$  different split-values for  $v$ . For each value of  $v$ , different machines need to compare the  $A_i$ -values of their portion of rows against  $v$ , and the resulting statistics need to be globally aggregated to compute the impurity of splitting. This cost is prohibitive since  $(n - 1)$  statistics transmissions are needed to examine just one attribute  $A_i$  for a node  $x$ . As a result, existing systems sacrifice accuracy for efficiency by only sampling a small set of splitting values of  $D_x$  for examination. For example, PLANET (adopted by MLlib and Mahout) computes approximate equi-depth histograms for each attribute [27], and a single splitting value is considered from every histogram bucket; XGBoost proposes a weighted quantile sketch for a similar purpose [17] but additionally updates sketches for each node  $x$  to reflect attribute values only in  $D_x$  rather than in the entire table.

PLANET also proposes to construct an entire subtree  $\Delta_x$  in a single machine when  $D_x$  is small enough, but such CPU-bound computation does not happen until towards the very end of tree construction, before which CPU cores remain underutilized. This is because PLANET constructs nodes top-down level by level, where upper-level nodes require inter-machine collaboration to find the best split-conditions (due to the row partitioning scheme of MapReduce upon which PLANET is built). PLANET constructs all nodes in a level by one MapReduce job, so that each row is read exactly once by a mapper task from HDFS to be IO-efficient. However, this design prevents tasks that build entire subtrees in a machine from being scheduled to run earlier to utilize idle CPU cores.

Among other related works, PV-Tree [23] proposes a communication-efficient heuristic to find split-condition ap-

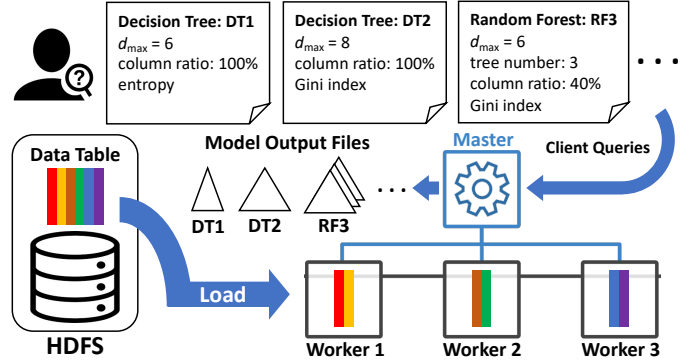


Fig. 2. Master-Workers Architecture of TreeServer

proximately in 3 steps: (1) selecting  $K$  attributes with the largest information gain from rows in each local machine, (2) finding  $2K$  global attributes among these candidates by majority voting, and (3) collecting histograms of these  $2K$  attributes to identify the best attribute and its split-condition. Yggdrasil [15] proposes to significantly reduce the data footprint by compression, and it finds the best split-conditions exactly. Their paper also analytically characterizes the impact of PLANET’s approximation to justify the need of exact computation. However, Yggdrasil still adopts a top-down level-by-level node construction order which is IO-bound and does not allow an entire subtree  $\Delta_x$  to be built in a single machine. Yggdrasil also uses a master to broadcast a bitvector of “row to child-node” assignment to all machines causing a single point of transmission bottleneck. Our TreeServer’s algorithm addresses both these problems. We remark that TreeServer can be extended to utilize compression similarly for further speedup, but as different attributes need different compression schemes which require additional user efforts to specify, we do not adopt it to keep our user API data-type transparent.

### III. SYSTEM OVERVIEW

**TreeServer Architecture.** Fig. 2 illustrates how TreeServer works. Specifically, TreeServer adopts a **master-workers** architecture. Users submit their model training jobs only to a master machine, which disassembles each tree model into individual decision trees for training. The master manages the training progress of each decision tree by distributing tasks to worker machines for computation and then collecting their results. The master then assembles the completed trees back into the target tree models to output. For example, in Fig. 2, a user submitted three jobs to train two decision trees DT1 and DT2, and a random forest RF3 with three trees, respectively. The master manages the training of all these decision trees (5 in total), and when completed, reassembles the trees into models DT1, DT2 and RF3 for output. Here, our master-workers architecture basically acts as a server for training individual decision trees (hence the name *TreeServer*).

Recall that TreeServer partitions the columns  $A_1, A_2, \dots, A_m$  of a data table among machines so that a machine holding  $A_i$  can compute the split-condition of  $A_i$  on its own. The columns are partitioned among the worker machines in a

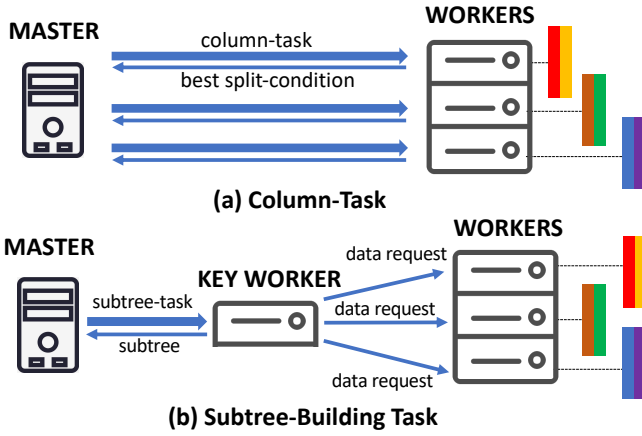


Fig. 3. Illustration of Subtree-tasks and Column-tasks

balanced manner: (1) the target column  $Y$  to predict is loaded into (the memory of) every machine since the impurity scores at node  $x$  are evaluated based on the  $Y$ -values of  $D_x$ ; (2) each other data column  $A_i$  is loaded into  $k$  machines. Fig. 2 shows the scenario where  $k = 1$ , but TreeServer uses  $k = 2$  by default since column replicas allow more room for load balancing: the task to compute the split-condition of a column  $A_i$  can be assigned to any machine that holds  $A_i$ , and we select the one with the minimum current workloads to avoid overloading busy machines; column replicas also support fault tolerance to avoid data loss due to a machine crash.

**Task Types.** In TreeServer, each task  $t_x$  is associated with a node  $x$  which builds the subtree  $\Delta_x$ . If  $D_x$  is large,  $t_x$  is split into subtasks associated with the child nodes of  $x$  to allow concurrent processing; while if  $D_x$  is small,  $t_x$  builds the entire  $\Delta_x$ . There are two possible task types for a task  $t_x$ :

- **Column-task.** To find the best split-condition of a node  $x$ , we need to first compute the best split-condition of each attribute  $A_i \in \mathcal{C}$ . When  $D_x$  is large, we let a machine that holds column  $A_i$  to compute the best split-condition of  $A_i$  for rows of  $D_x$ , which is called as a column-task. The best split-conditions computed by concurrent column-tasks are then gathered to find the overall best to split  $x$ .
- **Subtree-task.** If  $D_x$  is sufficiently small, we let a machine request attribute values of all the rows of  $D_x$  from other machines, which are then used to build subtree  $\Delta_x$ .

Fig. 3 illustrates these two types of tasks. Specifically, **if a task  $t_x$  has a large  $D_x$** , the master (1) distributes the column IDs to worker machines as *column-tasks* to find the best column-level split-conditions; it then (2) collects them to determine the overall best condition to split node  $x$  (see Fig. 3(a)); finally, the master (3) generates two subtasks  $t_{x_\ell}$  and  $t_{x_r}$  for  $x$ 's left child node  $x_\ell$  and right child node  $x_r$ , respectively, to build subtrees  $\Delta_{x_\ell}$  and  $\Delta_{x_r}$ . This essentially partitions the rows of  $D_x$  into  $D_{x_\ell}$  and  $D_{x_r}$  using the chosen split-condition. **On the other hand, if  $D_x$  is small**, the master lets a worker (called as the *key worker*) collect data of  $D_x$  from other machines to its local machine to build the subtree  $\Delta_x$ , which is then sent back to the master to be hooked to node  $x$  of the tree under construction (see Fig. 3(b)).

For now, let us assume each task  $t_x$  keeps  $I_x$  (i.e., the row indices of  $D_x$ ), but transmitting  $I_x$  during task assignment chokes the master's sending channel and is actually avoided in our design as shall be explained later in Section V.

**Tree Scheduling.** In T-thinker, only the master tracks the decision trees under construction; the other workers are only aware of column-tasks and subtree-tasks assigned by the master.

While our node-centric tasks allow parallel computation within each individual decision tree, in reality, we often need to train many tree models with different hyperparameters for model selection, or train ensemble models with many trees. T-thinker trains all these trees together so that we can have more node-centric tasks to keep CPUs busy. Individual decision trees and those in bagging (e.g., random forest, or a layer in deep forest) can be scheduled in parallel, while in boosting (e.g. gradient boosted trees, or layers in deep forest), sequential dependencies exist where the next layer of trees can only be scheduled for training when all trees in the previous layer is fully constructed. The master keeps track of the tree dependencies, and only considers a tree as a candidate for task scheduling if its prerequisite trees have all been constructed.

The master also controls the pace of tree construction when scheduling the tasks for eligible tree candidates: only a pool of  $n_{pool}$  trees are being actively constructed at any time. This keeps the memory usage bounded (as subtree-tasks gather  $D_x$  and thus consume memory) and prevents communication overload. When some trees are fully built and thus outputted to release the occupied memory, new eligible tree candidates can then have their tasks scheduled to use the available memory.

Note that no matter how many trees are specified to be trained by users (including those of ensemble models), a tree is flushed to disk by the master as soon as it receives the results from the tree's last task and completes the tree construction, which is tracked with the help of a progress table  $\mathcal{T}_{prog}$  at the master. See Appendix C for the implementation details.

**Task Scheduling.** Within each decision tree, upper-level nodes are processed by TreeServer in parallel using multiple machines to generate enough tasks as soon as possible to be assigned to the available CPU cores, while subtree-tasks of lower-level nodes are also scheduled timely to keep the CPU cores busy. This is achieved through a hybrid task scheduling scheme that combines breadth-first with depth-first traversal.

Specifically, TreeServer takes a job parameter  $\tau_D$  such that if  $|D_x| \leq \tau_D$ , task  $t_x$  is treated as a subtree-task (column-task otherwise). While subtree-tasks are CPU-bound, column-tasks incur IO workloads due to transmitting  $I_x$  (so that workers know which rows to check for columns  $A_i$  and  $Y$ ). TreeServer adopts a hybrid task scheduling scheme *combining breadth-first with depth-first node examination* to allow subtree-tasks and column-tasks to run concurrently to overlap computation with communication. When there are enough subtree-tasks, whenever a column-task is waiting for  $I_x$ , it is suspended so that its CPU core can be used to compute another task (e.g., a CPU-bound subtree-task). But when a TreeServer job begins, there are not enough subtree-tasks to saturate the available



CPU cores, so it is favorable to expand each tree level by level to generate more nodes  $x$  (hence tasks  $t_x$ ) to increase parallelism opportunities.

Fig. 4 illustrates the order that nodes (and also their tasks) in a tree are processed in our hybrid scheduling scheme, where upper-level nodes are processed in breadth-first order to quickly generate enough tasks for parallelism. Given a user-defined threshold  $\tau_{dfs}$  ( $> \tau_D$ ), when  $|D_x| \leq \tau_{dfs}$ , nodes in subtree  $\Delta_x$  are then processed in depth-first order (see the dashed subtree); moreover, during this depth-first node traversing, if  $|D_x| \leq \tau_D$ , then the entire  $\Delta_x$  is built by a subtree-task executed in one thread (see the black subtrees), in which case compute-bound subtree-tasks are timely scheduled.

Recall that TreeServer disassembles models to train as individual trees which are then trained concurrently, but only at most  $n_{pool}$  trees are under construction at any time. This design also helps to schedule subtree-tasks earlier to utilize CPU cores, since otherwise, we have to wait until all (and potentially many) tree “root” tasks have been expanded level by level to where depth-first node traversal begins. Now, when earlier trees in the pool are trained towards the end with subtree-tasks dominating, later comers are still at an early stage with column-tasks dominating, so both kinds of tasks are mixed.

To implement this task scheduler, the master maintains new tasks to be scheduled for computation in a **deque**  $B_{plan}$  as shown in Fig. 5 right, which is updated by 2 threads: (1) a main thread  $\theta_{main}$  that fetches tasks from the head of  $B_{plan}$  to be assigned for worker computation, and (2) a receiving thread  $\theta_{recv}$  that receives the results of column-tasks to find the best split-condition of each node  $x$  and that splits the node into subtasks of  $x$ ’s child-nodes to be inserted to  $B_{plan}$ . Specifically, when inserting a task  $t_x$ : (1) if  $t_x$  has  $|D_x| > \tau_{dfs}$ , then we append it to  $B_{plan}$  (like a queue) which essentially performs level-by-level breadth-first node examination, while if  $|D_x| \leq \tau_{dfs}$ , we add  $t_x$  to the head of  $B_{plan}$  (like a stack) which essentially performs depth-first examination of the nodes in  $\Delta_x$ .

To illustrate, consider Fig. 5 where a decision tree to build is shown on the left with  $|D_x|$  of each node  $x$  shown beside  $x$ . In Fig. 5 right, when the last column-task of node 3 is received by  $\theta_{recv}$  (see ①), master obtains the overall best split-condition




Fig. 4. Task Ordering




Fig. 5. An Example Decision Tree (Left) & Processing of Nodes 3 and 4 in the Master Machine (Right)

which splits rows of node 3 into child nodes 4 and 5 (see ②). As a result, child tasks for nodes 4 and 5 are created: node 4 is added to the head of  $B_{plan}$  since  $|D_4| \leq \tau_{dfs}$ , while node 5 is added to the tail of  $B_{plan}$  since  $|D_5| > \tau_{dfs}$  (see ③).

We tuned system parameters  $\tau_D$ ,  $\tau_{dfs}$  and  $n_{pool}$  and found the following default setting that works well consistently on various datasets:  $\tau_D = 10,000$ ,  $\tau_{dfs} = 80,000$ ,  $n_{pool} = 200$ . The experiments in Section VIII provide more details.

**Other TreeServer Features.** TreeServer also allows a tree path traversal to stop at any depth to output the predicted label at the current node, which provides additional flexibility in handling missing values and attribute values unseen during training. More details can be found in Appendix D.

#### IV. SYSTEM COMPONENTS

This section introduces the communication channels in T-thinker, and components of the master and workers. We also illustrate the workflow of task creation and processing.

##### Communication Channels.

As Fig. 6 shows, there are 2 types of communication channels: (1) “Task Comm.” for the master to send tasks to workers, and for workers to send task results back to the master; (2) “Data Comm.” among workers for data requesting and serving: recall that workers pull data from other workers to create  $D_x$  when they process subtree-tasks.


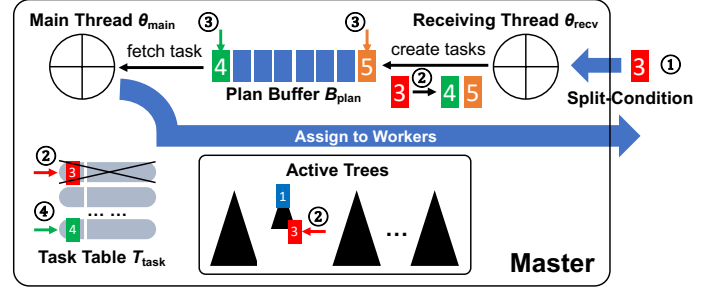


Fig. 6. Communication Channels

**Master Components.** The master is dedicated to task management and does not compute tasks by itself. Recall from Fig. 5 that the master has two threads: (1) main thread  $\theta_{main}$  for task assignment and (2) receiving thread  $\theta_{recv}$  that receives task results and updates the master status.



gement and does not compute tasks by itself. Recall from Fig. 5 that the master has two threads: (1) main thread  $\theta_{main}$  for task assignment and (2) receiving thread  $\theta_{recv}$  that receives task results and updates the master status.

In T-thinker, tasks are maintained by containers implemented with concurrent data structures for multi-threaded access, such as  $B_{plan}$  and  $T_{task}$  shown in Fig. 5. The task table  $T_{task}$  keeps the status for every task  $t_x$  in processing, so that it can be updated properly when  $\theta_{recv}$  receives the results related to  $t_x$ . For example, in Fig. 5, the task for node 3 (denoted by  $t_3$ ) was added to  $T_{task}$  after  $\theta_{main}$  fetches it from  $B_{plan}$  and assigns its column-tasks to the workers. Whenever  $\theta_{recv}$  receives the split-condition result w.r.t.  $t_3$  from a worker,  $\theta_{recv}$

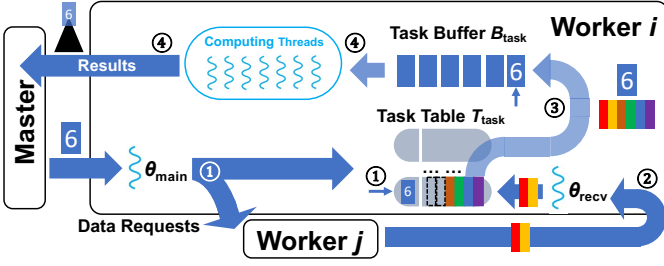


Fig. 7. Components of a Worker

uses it to update the current best split-condition of  $t_3$  in  $T_{table}$ , and updates  $t_3$ 's 'progress bar.' If the last split-condition result w.r.t.  $t_3$  is received by  $\theta_{recv}$ , the overall best split-condition is found, then node 3 is computed and thus added to its tree under construction, and  $t_3$ 's entry in  $T_{table}$  is freed.

Due to space limit, please refer to Appendix E for more details on task containers and thread workflows in the master.

### Worker Components.

Workers are the actual workhorses of task computation. Fig. 7 illustrates the components in a worker, where besides (1) a main thread  $\theta_{main}$  (see ①) for receiving tasks from the master and (2) a receiving thread  $\theta_{recv}$  (see ②) that receives data from other workers, the worker also maintains (3) a pool of computing threads (aka. **compers**) that fetch tasks from a task buffer  $B_{task}$  to compute their results (see ④), which are then sent back to the master.

Here,  $B_{task}$  keeps those tasks whose data are locally ready. If a subtree-task  $t_x$  is waiting for data from other works to construct  $D_x$ , then  $t_x$  is kept in the task table  $T_{task}$ , and only moved by  $\theta_{recv}$  to  $B_{task}$  when all required data are received.

Fig. 7 illustrates the execution flow of subtree-task  $t_6$  for node 6 (recall from Fig. 5 left that its  $|D_x| = 8,000 \leq \tau_D = 10,000$ ) in its key worker (recall from Fig. 3(b) that the key worker of a node  $x$  is the worker that collects  $D_x$  to build  $\Delta_x$ ). Specifically, when  $\theta_{main}$  receives the subtree task  $t_6$ , it sends requests for the columns of  $D_x$  to other workers which loaded these data columns, and meanwhile, it puts  $t_6$  to the task table  $T_{task}$  to wait for these columns to be collected (see ①). When  $\theta_{recv}$  receives the last few columns of  $D_x$  from another worker (see ②),  $D_x$  is fully collected so  $\theta_{recv}$  moves  $t_6$  out from  $T_{task}$  into  $B_{task}$  (see ③) to be fetched by an available compers to construct  $\Delta_x$ , which is then sent back to the master (see ④) to be added to the corresponding tree.

Due to space limit, please refer to Appendix E for more details on task containers and thread workflows in a worker.

**Fault Tolerance.** Since a TreeServer program is master-driven, master failure is fatal but can be addressed by keeping a

secondary master. If a worker crashes, the master simply reassigns the lost columns to other workers, revokes those tasks  $t_x$  that are being processed by the failed worker, and puts  $t_x$  back to  $B_{plan}$  for reassignment (recall  $B_{plan}$  from Fig. 5). More details on fault tolerance can be found in Appendix E.

### V. ROW MAINTENANCE WITHOUT MASTER RELAYING

So far, we have been assuming that each task (or task-plan) object  $t_x$  keeps the rows IDs of  $D_x$ , i.e.,  $I_x$ . As Fig. 8(a) illustrates, this applies to task-plans in  $B_{plan}$  and task objects in  $T_{task}$  in the master, as well as the assigned task-plans sent by the master to workers. However,  $I_x$  can be large for task-plans, leading to outbound communication bottleneck at the master. This section will present our technique to avoid letting any task (or task-plan) object  $t_x$  keep  $I_x$ , so that the communication bottleneck is eliminated. In the sequel, we first define some worker types related to our technique, and then present the motivation and details of our technique.

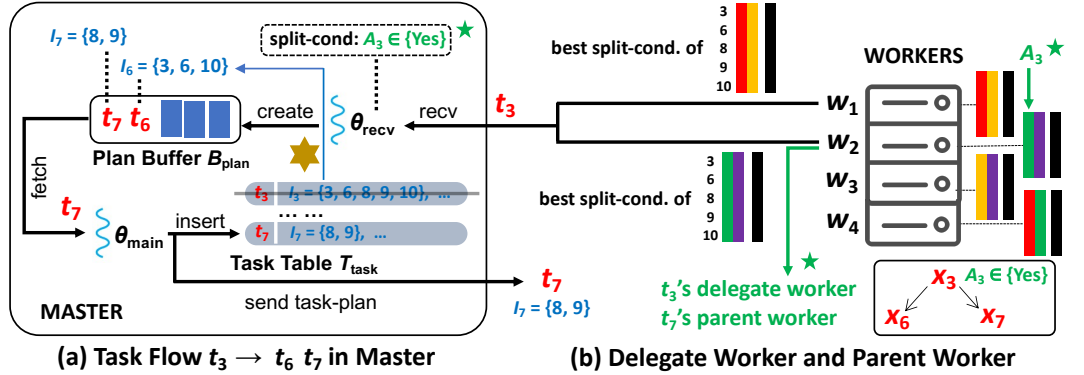


Fig. 8. Illustration of Delegate Worker and Parent Worker

**Worker Types.** Consider Fig. 1 again, where node  $x_3$  is split into nodes  $x_6$  and  $x_7$  using split-condition ' $A_3 \in \{\text{Yes}\}$ '. Fig. 8 shows this process, where we assume that the column-task  $t_3$  (i.e.,  $t_{x_3}$  but we abuse the notation for simplicity) generated two plans that were assigned to workers  $w_1$  and  $w_2$ , respectively. Also assume that the best split-condition from  $w_1$  as selected from its loaded columns  $A_1$  (red) and  $A_2$  (yellow) has been sent back to the master, as well as the best split-condition from  $w_2$  computed from its columns  $A_3$  (green) and  $A_4$  (purple). As a result,  $\theta_{recv}$  computes the overall best split-condition ' $A_3 \in \{\text{Yes}\}$ ' which gives the lowest impurity, and creates child tasks  $t_6$  and  $t_7$ .

One problem remains here: even though task object  $t_3 \in T_{task}$  maintains  $I_3$ , the master still does not have access to the attribute values of  $A_3$  to decide which new child node a row in  $D_x$  should go to. As a result, another round of communication would be necessary to send the computed split-condition ' $A_3 \in \{\text{Yes}\}$ ' to  $w_2$ , which splits  $I_3$  into  $I_6$  and  $I_7$  to be sent back to the master for adding to new child task-plans  $t_6$  and  $t_7$  to be added into  $B_{plan}$ .

In general, when a column-task  $t_x$  needs to generate child tasks  $t_{x_\ell}$  and  $t_{x_r}$  using the best split-condition obtained, assume that the best split-condition is with attribute  $A_i$  sent from Worker  $j$ , then in order to split  $I_x$  into  $I_{x_\ell}$  and  $I_{x_r}$ , we

have to let Worker  $j$  do the splitting since it has the attribute values of  $A_i$  to compare with the split-condition. Here, we call Worker  $j$  as the **delegate worker** of task  $t_x$ . For example, in Fig. 8,  $w_2$  that holds  $A_3$  is the delegate worker of task  $t_3$ .

Since the delegate worker splits  $I_x$  into  $I_{x_\ell}$  and  $I_{x_r}$ , ideally, we would like those workers that process  $t_{x_\ell}$  (resp.  $t_{x_r}$ ) to obtain  $I_{x_\ell}$  (resp.  $I_{x_r}$ ) directly from Worker  $j$ , rather than ask the master to relay  $I_{x_\ell}$  (resp.  $I_{x_r}$ ) in its task-plan message that creates repeated communication and outbound communication bottleneck. For example, in Fig. 8, when a task-plan for child task  $t_7$  is assigned to a worker, say  $w_4$ , we would like  $w_4$  to directly ask  $t_3$ 's delegate worker  $w_2$  for  $I_7$ , so that the master does not need to send  $I_7$  along with plan  $t_7$  (see Fig. 8(a)).

Here, we delegate the maintenance of  $I_x$  for a task  $t_x$  directly to its delegate worker, and task (or task-plan) objects  $t_x$  in  $B_{plan}$  and  $T_{task}$  also do not keep  $I_x$ . After all, in a constructed tree for prediction, each node  $x$  only needs to keep the split-condition rather than  $I_x$ , so the master does not have to track  $I_x$ . Since different tasks have different delegate workers, communication bottleneck at any machine is avoided.

Let us denote  $pa(x)$  as the parent node of  $x$ , and denote  $sib(x)$  as the sibling node of  $x$ . For example, in Fig. 8,  $pa(x_7) = x_3$  and  $sib(x_7) = x_6$ . Then, the worker that handles the requests for  $I_x$  is the delegate worker of the ‘‘parent’’ task  $t_{pa(x)}$ , which has the best split-attribute and thus can split  $I_{pa(x)}$  into  $I_x$  and  $I_{sib(x)}$  using the best split-condition confirmed by master, and which serves both child tasks  $t_x$  and  $t_{sib(x)}$ . We call the delegate worker of parent task  $t_{pa(x)}$  simply as the **parent worker** of task  $t_x$ .

**Technique Overview.** Fig. 9 shows the task workflows with our proposed technique, where the name ‘parent worker’ defined above is adopted. The delegate worker of a column-task  $t_x$  now cannot delete  $t_x$ 's task object from task table  $T_{task}$  immediately after the overall best split-condition is notified back by the master. This is because  $t_x$ 's task object holds  $I_x$  and hence  $I_{x_\ell}$  and  $I_{x_r}$  (after splitting  $I_x$ ), both of which will be requested by  $t_x$ 's child tasks later. After all the workers that process  $t_x$ 's child tasks have requested  $I_{x_\ell}$  and  $I_{x_r}$ ,  $t_x$ 's task object should then be deleted from  $T_{task}$ .

Also, since  $I_x$  is now tracked by  $t_x$ 's delegate worker rather than the master, we require a worker that sends back the best split-condition of a column  $A_i$  to also send two counters recording  $|I_{x_\ell}|$  and  $|I_{x_r}|$ , so that they can be compared with  $\tau_D$  to decide the type of child tasks (i.e., subtree-tasks or column-tasks) and whether to insert them into the head or tail of  $B_{plan}$ .

**Workflow of a Subtree-Task (Fig. 9(a)).** For a subtree-task  $t_x$ , the master assigns it to a worker for collecting  $D_x$  to build the subtree, and we call that worker as the **key worker** of  $t_x$ . The plan of a subtree-task indicates that for each attribute  $A_i \in \mathcal{C}$ , which machine should the key worker ask for attribute

values. The key-worker assignment and attribute-to-machine mapping (used by the key worker) are computed by the master.

As Fig. 9(a) shows, the master first sends a subtree plan  $t_x$  to its assigned key worker (see the arrow marked with ‘‘1’’), which creates a task object for  $t_x$  and puts it in task table  $T_{task}$  to collect  $D_x$ . The key worker then (1) requests other workers for attribute values of columns  $\mathcal{C}$  for rows of  $D_x$ , (2) sends a request to the parent worker to fetch  $I_x$  so that the key worker can obtain  $Y$ -values for those rows. These 2 kinds of requests are illustrated in Fig. 9(a) by the two arrows marked with ‘‘2’’ (where a thick arrow means multiple messages). When a worker receives the request for attribute values of column-subset  $\mathcal{C}' \subseteq \mathcal{C}$  (the request also contains the tracking information of  $t_{pa(x)}$  in the parent worker), it then creates a task object of  $t_x$  in task table  $T_{task}$  and requests  $I_x$  from the parent worker (see the arrow marked with ‘‘3’’). When the worker receives  $I_x$ , it then obtains  $\mathcal{C}'$  from  $t_x$ 's task object (the object is then deleted from  $T_{task}$ ), fetches these rows' attribute values in columns  $\mathcal{C}'$  and sends the data back to the key worker (see the arrow marked with ‘‘4’’). Finally, after the key worker collects the entire  $D_x$ , it builds subtree  $\Delta_x$  and sends it back to master (see the arrow marked with ‘‘5’’).

**Workflow of a Column-Task (Fig. 9(b)).** The master first sends the column-task plans of  $t_x$  to the assigned workers, each of which creates a task object of  $t_x$  (attached with the assigned column IDs  $\mathcal{C}'$  to find the best split-condition) and puts it in the task table  $T_{task}$  to wait for  $I_x$ , and then requests the parent worker for  $I_x$ . When  $I_x$  is received, the worker checks these rows for all columns in  $\mathcal{C}'$  to find the best split-




Fig. 9. Task Workflow with Delegate Worker

condition, which is then sent back to the master; by this time, the worker should not delete  $t_x$ 's task object yet. When the master has received all the responses of  $t_x$  and determines that Worker  $j$  has the best overall split-condition, it notifies all the other workers to delete their task objects for  $t_x$ . The task object of Worker  $j$  will be deleted by  $\theta_{recv}$  right after the last request from  $t_x$ 's child tasks is served, the progress of which is tracked by the task object and updated after each request for  $I_{x_\ell}$  or  $I_{x_r}$  (obtained by splitting  $I_x$  using the best split-condition) from a child task.

Assume that  $t_x$  is handled by  $k$  workers, then only  $t_x$ 's delegate worker needs to hold  $I_x$  for use by child tasks, while the other  $(k - 1)$  workers can delete their task object  $t_x$  immediately after node  $x$  is processed, leading to an additional  $1/k$  memory cost lingering till all child tasks have requested

$I_{x_\ell}$  or  $I_{x_r}$ . Since  $t_x$  is timely deleted by its delegate worker after being used, the additional memory cost is reasonable.

For readability, we omitted minor implementation details, such as skipping communication when the requested data is local, and the task workflow when the key worker of subtree-task  $t_x$  finds that  $D_x$  should be a leaf (e.g., when all  $Y$ -labels are the same so data request to collect  $D_x$  is not necessary).

## VI. WORKER ASSIGNMENT FOR TASKS

Since only the master handles task assignment, it has the big picture of plan-to-worker assignment and can keep track of the current worker workload distribution to assign each new plan to those workers that keep the workloads balanced.

Fig. 10 illustrates the matrix  $\mathcal{M}_{work}$  that the master uses to track workers' workloads, where each row corresponds to a worker and tracks its estimated current **computation**, **sending** and **receiving** workloads. For example, Row 1 indicates that Worker 1 needs to run 100 units of instructions for task computation, to send 70 units of messages to other machines, and to receive 60 units of messages, where the unit does not matter as long as they are the same for all workers, and the goal is to assign new plans to workers so that the values in every column of  $\mathcal{M}_{work}$  become as even (i.e., not drastically different) as possible.

Since (1) there are combinatorially many ways to assign task workloads of a new plan to workers (e.g., key worker and its data-serving workers), and (2) the workload dynamics of tasks that are currently being computed by workers (e.g., different steps in Fig. 9) are difficult to track including the delayed task object release at the delegate workers, it is impractical for the master to make decisions based on the accurate current worker workloads and consider all worker combinations.

To make plan assignment tractable and efficient, we adopt a greedy strategy that (1) treats all workloads of assigned tasks as pending even though they are partially processed, and that (2) makes decisions on each step based only on the dominant cost (i.e., computation or communication). When  $\theta_{main}$  of the master assigns the workloads of a task-plan  $t_x$  to workers, it updates  $\mathcal{M}_{work}$  by properly adding workloads; and when  $\theta_{recv}$  of the master receives the computation result for  $t_x$ , it deducts the added workloads (obtained from  $t_x$ 's task object in task table  $T_{task}$  which memorizes the added workloads) from  $\mathcal{M}_{work}$  to reflect the completion of  $t_x$ . A mutex protects  $\mathcal{M}_{work}$  so that only one thread ( $\theta_{main}$  or  $\theta_{recv}$ ) updates it at a time. We next describe our greedy plan-to-worker assignment.

**Assignment of a Subtree-Task.** Given a newly-created subtree-task plan  $t_x$ , we greedily assign the key worker as the worker with the minimum current **computation workload** in  $\mathcal{M}_{work}$  (since  $t_x$  is CPU-bound), and add its workload value by  $(|I_x| \cdot |\mathcal{C}| \cdot \log |I_x|)$ . Here, we assume that all attributes in  $\mathcal{C}$  are amenable to one-pass algorithm to find the best split-condition so that the incremental cost of attribute checking for

Worker ID	Comp	Send	Recv
1	100	70	60
2	60	80	75
3	80	70	90
4	120	60	85
...	...	...	...

Fig. 10. Load Matrix  $\mathcal{M}_{work}$

each row is  $O(|\mathcal{C}|)$  (as there are  $|\mathcal{C}|$  attribute values). Since the nodes of each level of  $\Delta_x$  partitions the rows of  $D_x$ , the overall checking cost for a level is  $|I_x| \cdot |\mathcal{C}|$ ; we also estimate the height of  $\Delta_x$  as  $\log |I_x|$  (assuming balanced  $\Delta_x$ ).

We next assign each column  $A_i \in \mathcal{C}$  to a worker so that the key worker of  $t_x$  will request it for  $A_i$ . We update  $\mathcal{M}_{work}$  after each  $A_i$  is assigned to a worker, so that the next column can be assigned by taking previous column assignments into consideration. Let us assume that the key worker is Worker  $i_{key}$  and the parent worker is Worker  $i_{pa}$ . If we assign column  $A_i$  to Worker  $j$ , then Worker  $j$  needs to (1) receive  $I_x$  from Worker  $i_{pa}$ , and to (2) send the  $|I_x|$  attribute values fetched from column  $A_i$  to Worker  $i_{key}$ . The two transmissions translate into the following 4 updates:

- 1)  $\mathcal{M}_{work}[j][Recv] \leftarrow \mathcal{M}_{work}[j][Recv] + |I_x|$
- 2)  $\mathcal{M}_{work}[i_{pa}][Send] \leftarrow \mathcal{M}_{work}[i_{pa}][Send] + |I_x|$
- 3)  $\mathcal{M}_{work}[j][Send] \leftarrow \mathcal{M}_{work}[j][Send] + |I_x|$
- 4)  $\mathcal{M}_{work}[i_{key}][Recv] \leftarrow \mathcal{M}_{work}[i_{key}][Recv] + |I_x|$

Among them, (3) and (4) are always needed; but to avoid double-counting, (1) and (2) are only needed if  $A_i \in \mathcal{C}$  is the first column assigned to Worker  $j$ . This is because Worker  $j$  will request Worker  $i_{pa}$  for  $I_x$  only once, for all columns  $\mathcal{C}'$  assigned to Worker  $j$  (including  $A_i$ ). To balance the network overheads among workers, Worker  $j$  is selected to minimize the maximum of the above 4 updated workload values.

**Assignment of a Column-Task.** Given a newly-created column-task plan  $t_x$ , we determine those workers that examine the columns  $\mathcal{C}$  to find the best split-condition. Note that each worker needs to request the parent worker for  $I_x$ , and then to examine the attribute values of its assigned columns  $\mathcal{C}'$  for these rows. Therefore, if we assign any column  $A_i$  to a worker  $j$ , we perform updates (1) and (2) exactly like before, and we also add  $\mathcal{M}_{work}[j][Comp]$  with the cost of examining  $A_i$  (e.g., the cost is  $|I_x|$  if a one-pass algorithm is used). Since the column assignment incurs network overheads as a major cost, Worker  $j$  is selected to minimize  $\max\{\mathcal{M}_{work}[j][Recv], \mathcal{M}_{work}[i_{pa}][Send]\}$  (after their values are updated) to balance communication.

TreeServer also properly skips adding communication workloads in special cases whenever the requested data are local.

## VII. A CASE STUDY WITH DEEP FOREST

Deep forest [34] is proposed as an alternative to deep neural networks, and it has reported an even higher accuracy. The model consists of many levels of forests each with hundreds of trees, making its training challenging for big data. This section explains how TreeServer can train a deep forest.

**Deep Forest Architecture.** Fig. 11 shows the structure of a deep forest, consisting of 2 phases: multi-grained scanning (MGS) and cascade forest (CF). MGS uses sliding windows of different sizes ( $\mathbf{W}_1, \mathbf{W}_2, \mathbf{W}_3$ ) to scan raw features, and the extracted window-sized vectors are used to train forests which are then used to re-represent raw features ( $\mathbf{f}_1, \mathbf{f}_2, \mathbf{f}_3$ ).

Fig. 12 illustrates the MGS process with a  $10 \times 10$  window, where the window slides on a raw image data with stride 1.



The extracted feature vectors have dimension  $10 \times 10 = 100$ , and such vectors from all images are used to train 2 forests. In deep forest, a forest for  $k$ -class classification returns a  $k$ -dimensional vector computed as the average of the class PMF vectors returned by all its trees. For each image, MGS concatenates the  $k$ -D vectors (outputted by both forests) for all window-sized vectors of the image as the inputs, and thus the re-representation can easily have thousands of dimensions.

As Fig. 11 shows, CF Layer 0 takes the re-represented data features from the smallest window as input, and trains forests which are then used to re-represent the data as the concatenation of  $k$ -D vectors outputted by the trained forests. In each later level, features outputted from the previous level are concatenated with the re-represented features from a window of particular size during MGS, to train forests to output the next-level re-representations. Prediction at each layer is obtained by averaging the  $k$ -D vectors outputted by the last-layer forests, then finding the mostly likely class.

**Implementation in TreeServer.** Since each CF layer trains multiple forests (each with many trees), it can be formulated as a TreeServer job. Similarly, training forests from window-sized features in MGS can be formulated as a TreeServer job. We also program 2 parallel operations to work with TreeServer jobs in an entire deep forest construction workflow:

- The first job is to perform window-sliding feature extraction over different images concurrently. For this purpose, image data are partitioned among the threads of all machines for concurrent window-sliding feature extraction.
- Once the forests of a layer have been constructed by TreeServer, they are saved to Hadoop Distributed File System (HDFS). Then, the image data needs to go through these forests again (i.e., the prediction phase) to generate output features (aka. re-representations), which is the task of our second job. Here, we let every machine load all the forests from HDFS, and then conduct tree traversal for its assigned portion of images. Our second job allows different images to be concurrently processed by different threads of different machines.

These two operations partitions input data by rows, while TreeServer model training partitions input data by columns. Their integration calls for a novel data organization on HDFS to be described soon that is friendly to both partition schemes.




Fig. 12. Image Feature Re-representation

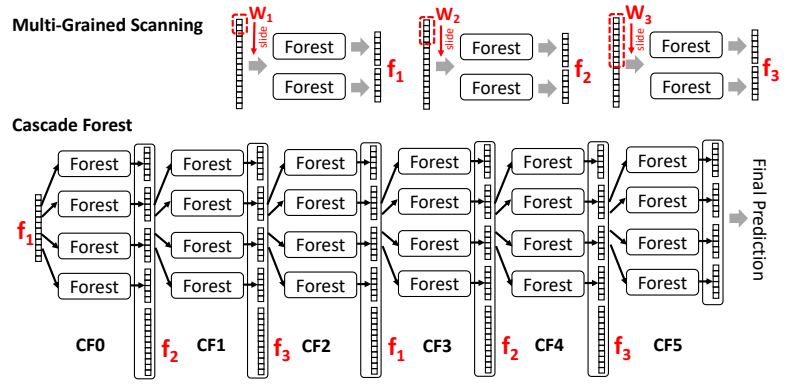


Fig. 11. Overall Procedure of Deep Forest

Besides random forests, deep forest also supports “completely random decision trees” (aka. extra-trees) as the other forest type. See Appendix F for the details.

**Data Organization on HDFS.** To support TreeServer’s column partitioning scheme, we require users to use our dedicated “put” program (instead of HDFS’s) to upload data (e.g., a local CSV file) to HDFS, so that each data column is saved as a file on HDFS that can be loaded by workers in its entirety. Our “put” program is memory-efficient: it uses  $m$  HDFS output file streams (one for each column) to append the corresponding attribute values in each data row, as it streams data rows.

A normal tabular data has at most tens to several hundreds of columns, but as we have explained, MGS can easily generate data representations of thousands of dimensions or even more. The many column files would be inefficient for workers to load since HDFS connection time (rather than actual data reads) dominates in our test. Our solution is to group columns to reduce the number of files, as illustrated in Fig. 13 where the input data is organized into 4 column-files each with 50 attributes. Now the column-groups (rather than individual columns) are assigned as the basic unit to the workers in a balanced manner to serve data requests by tasks.

Another necessary change is due to our two additional parallel programs described earlier that partitions data by rows. To support both row- and column-partitioning efficiently, our final data organization adopted is as illustrated in Fig. 13, where each column-group is further partitioned into row-groups, and one file is saved for each row-group. In Fig. 13, now a TreeServer job may load a column-group by reading 4 files in the same column, while our two new parallel programs may load its partition of rows by reading 4 files in the same row. Here, the number of files to read is kept small, and each file contains sufficient data to amortize HDFS connection cost.

columns	$A_1 - A_{50}$	$A_{51} - A_{100}$	$A_{101} - A_{150}$	$A_{151} - A_{200}$
rows				
1 - 250				
251 - 500				
501 - 750				
751 - 1000				

Fig. 13. Data Organization on HDFS

## VIII. EXPERIMENTS

This section evaluates the efficiency of TreeServer, and compares it with the tree models in Spark MLlib [24] and XGBoost [17], the state-of-the-art machine libraries for tree models that support distributed training. Our code is released on GitHub at: <https://github.com/yanlab19870714/TreeServer>.

All our experiments were conducted on a cluster of 15 machines each with 12 threads (2.67 GHz) and 24 GB RAM. Each experiment was repeated 3 times, with the average results reported. Unless otherwise stated, we run with all 15 machines each with 10 compers, and we set the maximum tree depth  $d_{max} = 10$ ,  $\tau_{leaf} = 1$ , classification (resp. regression) impurity as Gini index (resp. variance), and  $|\mathcal{C}| = |\mathcal{A}|$  for decision trees and  $|\mathcal{C}| = \sqrt{|\mathcal{A}|}$  for a tree in random forests.

TABLE I  
DATASETS

Dataset	#{rows}	#{numerical}	#{categorical}	Problem
Allstate	13,184,290	13	14	regression
Higgs_boson	11,000,000	28	0	classification
MS_LTRC	723,412	136	1	
c14B	473,134	700	0	
Covtype	581,012	54	0	
Poker	1,025,010	0	11	
KDD99	4,898,431	38	3	
SUSY	5,000,000	18	0	
loan_m1	6,372,703	14	13	
loan_y1	29,581,722			
loan_y2	54,468,375			

**Datasets.** Table I shows the large real datasets we used, each of which either has many rows or many attributes. For all the datasets, we have removed attributes that are not related to prediction, such as ID, date and time, etc. *Allstate* [1], *Higgs\_boson* [3], *MS\_LTRC* [5], *c14B* [14], *Covtype* [2], *Poker* [6], *KDD99* [4] and *SUSY* [8] are directly usable, among which only *Allstate* is a regression problem and has missing values. TreeServer and XGBoost handle missing values and can directly take *Allstate* as input; while Spark MLlib does not support missing values and so we had to fill missing values with the mean attribute value first. The other dataset *loan* [7] requires preprocessing (see Appendix G for details), and we took the loan data from the latest month, the latest year, and the latest 2 years to obtain 3 datasets of different sizes.

**Comparison with MLlib.** Spark MLlib uses an argument “maxBins” to indicate the binning size of its attribute value histogram used for finding approximate split-conditions, and we use the default maxBins = 32. Table II(a) (resp. Table II(b)) shows the training time of both TreeServer and MLlib on our datasets for training a decision tree using all columns (resp. training a random forest with 20 trees using  $\sqrt{|\mathcal{A}|}$  columns). We can see that TreeServer is consistently many times faster than MLlib (e.g., almost 10 $\times$  on *MS\_LTRC*). Moreover, TreeServer computes the exact split-conditions while MLlib computes split-conditions approximately with attribute binning, and as shown in Tables II(a) & (b), the test accuracy of models trained with TreeServer is slightly higher than that of models from MLlib in the majority of the cases.

Table II(c) compares TreeServer with XGBoost both with 100 trees. Note that TreeServer uses bagging (i.e., random for-

TABLE II  
SYSTEM COMPARISON (ACCURACY = RMSE FOR ALLSTATE)

(a) TreeServer v.s. MLlib One Decision Tree

Dataset	TreeServer		MLlib (Parallel)		MLlib (Single Thread)	
	Time (sec)	Accuracy	Time (sec)	Accuracy	Time (sec)	Accuracy
Allstate	56.06	41.46	194.26	41.26	886.98	41.26
Higgs_boson	103.53	70.52%	137.36	69.16%	467.38	69.16%
MS_LTRC	12.21	55.30%	103.45	55.38%	14.60	55.38%
c14B	37.81	67.00%	119.49	49.04%	252.61	49.04%
Covtype	10.00	95.31%	58.39	95.29%	67.14	95.29%
Poker	10.08	53.52%	86.55	53.52%	85.45	53.52%
KDD99	58.58	79.98%	83.90	79.97%	243.00	79.97%
SUSY	38.68	79.13%	90.76	79.10%	156.35	79.10%
loan_m1	52.49	99.63%	322.47	99.50%	524.17	99.50%
loan_y1	284.75	99.64%	551.64	99.64%	1,188.26	99.64%
loan_y2	520.68	99.64%	808.09	99.64%	1,972.12	99.64%

(b) TreeServer v.s. MLlib Random Forest (20 Trees)

Dataset	TreeServer		MLlib (Parallel)		MLlib (Single Thread)	
	Time (sec)	Accuracy	Time (sec)	Accuracy	Time (sec)	Accuracy
Allstate	113.31	41.46	305.86	41.39	1,635.70	41.39
Higgs_boson	135.53	69.89%	259.75	69.69%	978.69	69.69%
MS_LTRC	14.43	55.23%	122.47	55.59%	14.67	55.59%
c14B	29.86	62.00%	125.01	49.59%	277.67	49.59%
Covtype	10.82	91.53%	82.38	91.52%	79.42	91.52%
Poker	10.06	53.71%	174.15	53.95%	175.81	53.95%
KDD99	46.59	80.15%	234.67	80.15%	414.85	80.15%
SUSY	51.81	79.49%	127.59	79.37%	287.29	79.37%
loan_m1	42.44	99.64%	373.25	99.50%	560.28	99.50%
loan_y1	192.73	99.64%	786.94	99.64%	2,856.16	99.64%
loan_y2	367.84	99.64%	1422.20	99.64%	5,008.66	99.64%

(c) TreeServer v.s. XGBoost (100 Trees)

Dataset	TreeServer		XGBoost	
	Time (sec)	Accuracy	Time (sec)	Accuracy
Allstate	381.94	41.46	757.77	38.81
Higgs_boson	429.82	69.85%	2,364.87	70.12%
MS_LTRC	45.38	55.22%	1,585.10	58.83%
c14B	62.33	61.00%	3,492.70	55.85%
Covtype	47.66	92.33%	767.78	93.67%
Poker	23.06	55.95%	1,338.56	56.87%
KDD99	186.03	99.91%	3,892.76	99.93%
SUSY	141.92	79.88%	1,650.33	79.54%
loan_m1	98.52	99.64%	1405.08	99.47%
loan_y1	540.88	99.64%	7,411.35	99.47%
loan_y2	788.66	99.64%	13,390.34	99.47%

est) while XGBoost uses boosting. We can see from Table II(c) that XGBoost achieves a higher accuracy on 6 out of the 11 datasets, thanks to its advanced gradient boosting method that considers second-order approximation of the learning objective, but not higher on the other datasets. In contrast, it is much slower than TreeServer. For example, on *C14B* (resp. *Poker*), XGBoost is 56 $\times$  (resp. 58 $\times$ ) slower than TreeServer. This is because trees in XGBoost have dependencies and have to be trained one after another due to boosting, while trees in TreeServer can be trained together without any dependency.

**Effect of  $n_{pool}$ .** Recall that to control memory usage, we only allow  $n_{pool}$  trees to be built at any time, whose value is 200 by default. To explore the relationship between  $n_{pool}$  and a job’s peak memory usage (at a machine, averaged over all machines) and running time, we trained a random forest with 20 trees on the various datasets with  $n_{pool} = 1, 5, 10$  and 20, respectively. Tables III(a)–(c) show our results on *Allstate*, *Higgs\_boson*, and *KDD99* (results on other datasets are similar). We can see that the job running time decreases a lot as  $n_{pool}$  increases, though the improvement is less significant as  $n_{pool}$  gets larger (e.g., 6 $\times$  on *Allstate* when  $n_{pool} = 20$ ). This is because our machine is equipped with a Xeon X5650 CPU (6 cores, 12 threads) which is already near saturation given the tasks of



TABLE VI  
HORIZONTAL SCALABILITY (TIME UNIT: SECOND; NET UNIT: MBPS)

(a) TreeServer (1 Tree)							(b) Spark MLlib (1 Tree)						
#(macs)	Allstate			Higgs_boson			#(macs)	Allstate			Higgs_boson		
	Time	CPU	Send	Time	CPU	Send		Time	CPU	Send	Time	CPU	Send
4	123.76	837%	605.61	174.64	860%	940.40	4	241.77	167.99		241.77	167.99	
8	76.72	760%	916.41	107.16	881%	942.32	8	244.13	180.93		244.13	180.93	
12	65.74	699%	941.26	103.48	800%	942.52	12	195.78	134.45		195.78	134.45	
15	61.93	694%	941.34	120.91	817%	942.32	15	195.99	139.52		195.99	139.52	

(c) TreeServer (20 Tree)							(d) Spark MLlib (20 Tree)						
#(macs)	Allstate			Higgs_boson			#(macs)	Allstate			Higgs_boson		
	Time	CPU	Send	Time	CPU	Send		Time	CPU	Send	Time	CPU	Send
4	236.79	726%	893.49	314.28	885%	941.33	4	366.51	355.69		366.51	355.69	
8	146.26	710%	935.45	193.10	840%	941.48	8	388.22	340.00		388.22	340.00	
12	142.78	716%	941.48	146.13	810%	941.28	12	304.52	263.09		304.52	263.09	
15	118.18	663%	941.28	145.73	815%	942.25	15	308.25	259.75		308.25	259.75	

TABLE VII  
DEEP FOREST EXPERIMENTAL RESULTS

Cascade Forest			Step	Training Time	Test Time	Test Accuracy
Multi-Grained Scanning			CF0train	94.88 s	—	—
			CF0extract	41.04 s	27.07 s	96.87%
			CF1train	99.31 s	—	—
			CF1extract	41.95 s	28.08 s	96.82%
			CF2train	93.55 s	—	—
			CF2extract	39.63 s	27.44 s	97.68%
			CF3train	94.73 s	—	—
			CF3extract	37.68 s	26.84 s	98.10%
			CF4train	91.61 s	—	—
			CF4extract	41.99 s	27.93 s	97.02%
Step	Training Time	Test Time	CF5train	92.46 s	—	—
slide	163.58 s	34.95 s	CF5extract	40.07 s	27.63 s	97.71%
win3train	716.57 s	—				
win5train	971.93 s	—				
win7train	1251.33 s	—				
win3extract	105.00 s	38.69 s				
win5extract	127.72 s	45.06 s				
win7extract	117.85 s	42.57 s				

running 10 compers, we can achieve 837% average CPU rate (meaning that 8 cores are fully utilized). As we increase machine number to 15, CPU rate is still 694% but since more data transmission is required, average sending throughput saturates at 941.34 Mbps (we use 1 GigE) so performance improvement flattens. We expect the performance to improve further with more machines if 10 GigE is available. As Tables II(a) and (c) illustrates, we are able to consistently keep  $> 6$  (oftentimes  $> 8$ ) CPU cores busy while fully utilizing the available network bandwidth. In contrast, MLlib sees less significant improvement and is more expensive in all cases.

**Experiments on Deep Forests.** We run our deep forest pipeline built with TreeServer on the MNIST dataset which consists of  $28 \times 28$  gray images, 60,000 (resp. 10,000) of them are for training (resp. testing). All images are for digits 0–9, ten classes in total. We used MGS window size  $3 \times 3$ ,  $5 \times 5$  and  $7 \times 7$ , and we followed Table 1 of [34] to set the hyperparameters but make the following changes to improve accuracy: (1) extra trees reduce the test accuracy when used in CF so we only use random forests for CF, (2) setting  $d_{max} = 100$  in MGS drops the test accuracy so we use 10, (3) to reach good accuracy, each step only needs 2 forests each with 20 trees (rather than 500 trees), and only 10% of the training and test images are needed and thus used.

Table VII reports the running time of each step as illustrated in Fig. 11, as well as their corresponding test accuracy. Specifically, “slide” corresponds to the step that slides each image to generate inputs to MGS. In MGS, “win3train”, for example, corresponds to the step that trains the forests over vectors extracted by  $3 \times 3$  windows, and “win3extract”

TABLE VIII  
IMPACT OF MODEL PARAMETERS

(a) Effect of $d_{max}$ on Higgs_boson (1 tree)			(b) Effect of $d_{max}$ on Higgs_boson (20 trees)		
$d_{max}$	Time	Accuracy	$d_{max}$	Time	Accuracy
2	69.47 s	63.117%	2	84.65 s	62.372%
4	93.37 s	65.635%	4	115.99 s	61.571%
6	118.62 s	67.946%	6	124.21 s	62.655%
8	112.96 s	69.456%	8	128.73 s	63.546%
10	115.86 s	70.515%	10	155.68 s	63.893%
12	113.04 s	71.313%	12	170.11 s	64.153%

(c) Effect of $ C / A $ on Allstate			(d) Effect of $ C / A $ on Higgs_boson		
$ C / A $	Time	RMSE	$ C / A $	Time	Accuracy
20%	121.79 s	41.460	20%	158.66 s	63.893%
40%	145.71 s	41.469	40%	266.23 s	70.355%
60%	220.35 s	41.462	60%	416.07 s	71.038%
80%	241.02 s	41.463	80%	423.70 s	71.506%
100%	255.14 s	41.464	100%	622.49 s	70.515%

corresponds to the step that used the trained forests to represent each input image. In the CF stage, “CF0train”, for example, corresponds to the time to train the forests of CF0 while “CF0extract” corresponds to the time to use the trained forest in CF0 to represent the input images. Note that after each CF step, we report our test accuracy which are all very high and improves all the way up to CF3. We observe that training accuracy is always 100% because the maximum tree depth in the CF stage is  $d_{max} = \infty$ , so we omit them in Table VII. We can see that training is very efficient despite the many trees.

**Fairness of Implementation.** Since MLlib runs on JVM, while our TreeServer program runs as a native C++ program, it is worth exploring whether our improvement comes from the language rather than the system design. For this purpose, we run single-threaded single-tree construction for both TreeServer and MLlib and find that TreeServer is comparable to MLlib. For example, TreeServer takes 705.94 seconds while MLlib takes 750.58 seconds on *Higgs\_boson*, and TreeServer takes 191.86 seconds while MLlib takes 157.34 seconds on *MS\_LTRC*. Our single-tree implementation is not faster since in order to allow flexible user data input like in pandas, we conduct runtime (rather than compile-time) type conversion based on the data type of each column being processed.

**Accuracy w.r.t. Model Parameters.** Since TreeServer implements the exact node splitting algorithms, we obtain exact models as in a conventional serial algorithm but runs much faster. Table VIII(a) (resp. (b)) shows the time and test accuracy of training a tree (resp. forest with 20 trees) on *Higgs\_boson* with increasing  $d_{max}$  and we can see that the accuracy keeps improving (i.e., models are not overfitting). Tables VIII(c)–(d) show the performance of training a 20-tree forest with different  $|C|$  sampled for each tree on *Allstate* and *Higgs\_boson*, where the test accuracy does not change a lot meaning that 20% columns per tree are already sufficient.

## IX. CONCLUSION

We presented a distributed system called TreeServer for training tree models, which addresses the problems of IO bottleneck in existing solutions such as MLlib. The system adopts novel designs including column-based data partitioning and node-centric task-based workload partitioning, and is consistently many times faster than MLlib, hence a good alternative especially for training big models like a deep forest.



## REFERENCES

- [1] Allstate Dataset. <https://www.kaggle.com/c/allstate-claims-severity/data>.
- [2] Forest CoverType Dataset. <https://archive.ics.uci.edu/ml/datasets/covertype>.
- [3] Higgs Boson Dataset. <https://archive.ics.uci.edu/ml/datasets/HIGGS>.
- [4] KDD Cup 1999 Dataset. <https://archive.ics.uci.edu/ml/datasets/kdd+cup+1999+data>.
- [5] Microsoft Learn2Rank Dataset. <https://www.microsoft.com/en-us/research/project/mslr/>.
- [6] Poker Hand Dataset. <https://archive.ics.uci.edu/ml/datasets/poker+hand>.
- [7] Single Family Loan-Level Dataset. [http://www.freddiemac.com/research/datasets/sf\\_loanlevel\\_dataset.page](http://www.freddiemac.com/research/datasets/sf_loanlevel_dataset.page).
- [8] Supersymmetry (SUSY) Dataset. <https://archive.ics.uci.edu/ml/datasets/SUSY>.
- [9] TreeServer Demo (Long Version). [https://www.youtube.com/watch?v=4DnLv\\_OFIIg](https://www.youtube.com/watch?v=4DnLv_OFIIg).
- [10] TreeServer Demo (Short Version). <https://www.youtube.com/watch?v=IawT40DhGbs>.
- [11] TreeServer GitHub Repository. <https://github.com/yanlab19870714/TreeServer>.
- [12] TreeServer's Deep Forest Demo. [https://www.youtube.com/watch?v=LzXkzkk0r\\_0](https://www.youtube.com/watch?v=LzXkzkk0r_0).
- [13] TreeServer's Deep Forest GitHub Repo. <https://github.com/yanlab19870714/deepForest>.
- [14] Yahoo Learn2Rank Dataset. <https://webscope.sandbox.yahoo.com/catalog.php?datatype=c>.
- [15] F. Abuzaid, J. K. Bradley, F. T. Liang, A. Feng, L. Yang, M. Zaharia, and A. S. Talwalkar. Yggdrasil: An optimized system for training deep decision trees at scale. In *NIPS*, pages 3810–3818, 2016.
- [16] L. Breiman, J. H. Friedman, R. A. Olshen, and C. J. Stone. *Classification and Regression Trees*. Wadsworth, 1984.
- [17] T. Chen and C. Guestrin. Xgboost: A scalable tree boosting system. In *SIGKDD*, pages 785–794, 2016.
- [18] W. Fan, H. Wang, P. S. Yu, and S. Ma. Is random model better? on its accuracy and efficiency. In *ICDM*, pages 51–58, 2003.
- [19] J. Feng, Y. Yu, and Z. Zhou. Multi-layered gradient boosting decision trees. In *NeurIPS*, pages 3555–3565, 2018.
- [20] Frank McSherry. COST in the land of databases. <https://github.com/frankmcsherry/blog/blob/master/posts/2017-09-23.md>, 2017.
- [21] G. Guo, D. Yan, M. T. Özsu, and Z. Jiang. Scalable mining of maximal quasi-cliques: An algorithm-system codesign approach. *PVLDB*, 2021.
- [22] J. Khalil, D. Yan, G. Guo, and L. Yuan. Parallel mining of large maximal quasi-cliques. *The VLDB Journal (to appear)*, 2021.
- [23] Q. Meng, G. Ke, T. Wang, W. Chen, Q. Ye, Z. Ma, and T. Liu. A communication-efficient parallel algorithm for decision tree. In *NIPS*, pages 1271–1279, 2016.
- [24] X. Meng, J. K. Bradley, B. Yavuz, E. R. Sparks, S. Venkataraman, D. Liu, J. Freeman, D. B. Tsai, M. Amde, S. Owen, D. Xin, R. Xin, M. J. Franklin, R. Zadeh, M. Zaharia, and A. Talwalkar. Millib: Machine learning in apache spark. *J. Mach. Learn. Res.*, 17:34:1–34:7, 2016.
- [25] M. M. Michael and M. L. Scott. Simple, fast, and practical non-blocking and blocking concurrent queue algorithms. In *Proceedings of the Fifteenth Annual ACM Symposium on Principles of Distributed Computing, Philadelphia, Pennsylvania, USA, May 23-26, 1996*, pages 267–275, 1996.
- [26] A. Musselman. Apache mahout. In *Encyclopedia of Big Data Technologies*. 2019.
- [27] B. Panda, J. Herbach, S. Basu, and R. J. Bayardo. PLANET: massively parallel learning of tree ensembles with mapreduce. *PVLDB*, 2(2):1426–1437, 2009.
- [28] W. Qu, D. Yan, G. Guo, X. Wang, L. Zou, and Y. Zhou. Parallel mining of frequent subtree patterns. In *LSGDA@VLDB*, 2020.
- [29] D. Yan, G. Guo, M. M. R. Chowdhury, M. T. Özsu, W.-S. Ku, and J. C. S. Lui. G-thinker: A distributed framework for mining subgraphs in a big graph. In *ICDE*, 2020.
- [30] D. Yan, G. Guo, M. M. R. Chowdhury, M. T. Özsu, J. C. S. Lui, and W. Tan. T-thinker: a task-centric distributed framework for compute-intensive divide-and-conquer algorithms. In *PPoPP*, pages 411–412, 2019.
- [31] D. Yan, G. Guo, J. Khalil, M. T. Özsu, W.-S. Ku, and J. Lui. G-thinker: a general distributed framework for finding qualified subgraphs in a big graph with load balancing. *The VLDB Journal*, pages 1–34, 2021.
- [32] D. Yan, W. Qu, G. Guo, and X. Wang. Prefixfp: A parallel framework for general-purpose frequent pattern mining. In *ICDE*, 2020.
- [33] D. Yan, W. Qu, G. Guo, X. Wang, and Y. Zhou. Prefixfp: a parallel framework for general-purpose mining of frequent and closed patterns. *The VLDB Journal*, pages 1–34, 2021.
- [34] Z. Zhou and J. Feng. Deep forest: Towards an alternative to deep neural networks. In *IJCAI*, pages 3553–3559, 2017.

### A. List of Notations

The paper used the following notations, which are summarized for quick reference as needed:

- $D$ : the input data table;
- $n$ : number of rows in the data table;
- $m$ : number of attributes (i.e., columns) in the data table;
- $A_i$ : the  $i$ -th attribute in the data table;
- $\mathcal{A}$ : the set of all attributes  $A_1, A_2, \dots, A_m$ ;
- $Y$ : the target attribute to predict;
- $x$ : a node in a decision tree;
- $\Delta_x$ : the subtree rooted at node  $x$ ;
- $D_x$ : those rows of the training data that reach node  $x$  following the decision path from the root node;
- $|D_x|$ : the number of rows in  $D_x$ ;
- $I_x$ : the row-IDs of those rows in  $D_x$ ;
- $v$ : the split-value of a numerical attribute  $A_i$ , used in split-condition “ $A_i \leq v$ ”;
- $S_i$ : the set of all possible values of a categorical attribute  $A_i$ ;
- $S_\ell$ : a subset of attribute values in  $S_i$  which lead a data to go to the left child node of  $x$ , used in split-condition “ $A_i \in S_\ell$ ”;
- $\mathcal{C}$ : a subset of  $\mathcal{A}$ , which restricts the split-condition of node  $x$  to be chosen only from an attribute in  $\mathcal{C}$ ;
- $|\mathcal{C}|$ : a number of columns in  $\mathcal{C}$ ;
- $t_x$ : the task that processes node  $x$  to construct its subtree;
- $x_\ell$ : the left child node of node  $x$ ;
- $x_r$ : the right child node of node  $x$ ;
- $\tau_D$ : a system threshold such that if  $|D_x| \leq \tau_D$ , the entire subtree  $\Delta_x$  is built by a subtree-task  $t_x$ ;
- $\tau_{dfs}$ : a system threshold such that if  $|D_x| \leq \tau_{dfs}$ , nodes in the subtree  $\Delta_x$  are processed in depth-first order (breadth-first order otherwise), and usually  $\tau_{dfs} > \tau_D$  (see Fig. 4);
- $n_{pool}$ : the maximum number of trees allowed to be under construction at any time;
- $B_{plan}$ : a buffer (deque) in the master to hold tasks (aka. plans) newly created but not assigned worker workloads yet;
- $d_{max}$ : the maximum allowed tree depth, which is a model hyperparameter;
- $\tau_{leaf}$ : a model hyperparameter such that a node  $x$  stops its splitting to avoid overfitting when  $|D_x| \leq \tau_{leaf}$ ;
- $Q_{plan}$ : a plan queue in the master to buffer task-plans that have been assigned worker workloads, and these plans will be sent to the destination workers in batches;
- $T_{task}$ : a task table to keep tasks that are being processed but are waiting for dependent data;
- $\theta_{main}$ : the main thread in a machine;
- $\theta_{recv}$ : a communication thread that receives responses from other machines;
- $comper$ : a computing thread in a worker machine;
- $B_{task}$ : a task buffer of a worker machine from which compers take tasks for computation;

- $\mathcal{M}_{work}$ : a workload matrix that the master uses to track worker workloads, where each row  $i$  is for a worker, and  $\mathcal{M}_{work}[i][Comp]$ ,  $\mathcal{M}_{work}[i][Send]$  and  $\mathcal{M}_{work}[i][Recv]$  tracks Worker  $i$ ’s computation, sending and receiving workloads, respectively.

### B. Selecting Best Split-Condition for a Table Column

In decision tree training, at each tree node  $x$ , we select the best split-condition for each individual attribute (column)  $A_i$  independently. Depending on the type of  $A_i$  and  $Y$ , there are three cases:

- **Case 1:  $A_i$  is an ordinal attribute.** To find the best value  $v$  for the split-condition “ $A_i \leq v$ ”, only the values of  $A_i$  that appear in  $D_x$  are considered. It is well known that by sorting rows of  $D_x$  by  $A_i$ -values and checking each value  $v$  to split the rows into left and right child nodes, we can get the impurity value for each  $v$  in  $O(1)$  incremental cost and thus we can find the best  $v$  in one pass over the sorted  $D_x$ .
- **Case 2:  $A_i$  is categorical, and  $Y$  is numerical.** This case is for regression, and Breiman et al. [16] present an algorithm for finding the best split predicate without evaluating all possible subsets of  $S_i$  for  $S_\ell$ . After grouping rows of  $D_x$  by  $A_i$ -values and sorting the groups by the average  $Y$ -value, the optimal split predicate cuts the sorted group list in the middle, so the algorithm only needs one pass over the groups.
- **Case 3: both  $A_i$  and  $Y$  are categorical.** This case is for classification, and we have to enumerate and check all possible subsets of  $S_i$  for  $S_\ell$  to split rows in  $D_x$ . To limit the checking cost when  $|S_i|$  is big, it is common to restrict  $|S_\ell| = 1$  so that only  $O(|S_i|)$  split-conditions need to be checked.

### C. Tree Construction Progress Tracking

The master uses a progress table  $\mathcal{T}_{prog}$  to track the number of pending tasks for each active tree under construction, so that once it receives the results from the tree’s last task and completes the tree construction, it can reduce the count of active trees so that new trees can be admitted for processing (recall that TreeServer executes a budget of at most  $n_{pool}$  active trees at any moment).

In our implementation, each task keeps its tree ID  $tid$ . Here, a tree “root” task is created by setting its  $tid$  properly when the tree is admitted for processing, while each descendent node-centric task  $t_x$  inherits its tree ID  $tid$  from its parent task  $t_{pa(x)}$ , where  $pa(x)$  is the parent node of node  $x$ .

Each task sends its tree ID back to the master along with the task result, to update the progress counter  $\mathcal{T}_{prog}[tid]$ : incremented if the task is split into two child tasks, and decremented otherwise. This is because a column-task is consumed by the master which creates two new child tasks, leading to a net increment of 1; while a subtree-task is consumed without creating new tasks, leading to a net increment of -1.

As soon as the master finds that  $\mathcal{T}_{prog}[tid]$  is decremented to 0 after processing the result of a task from Tree  $tid$ , it knows

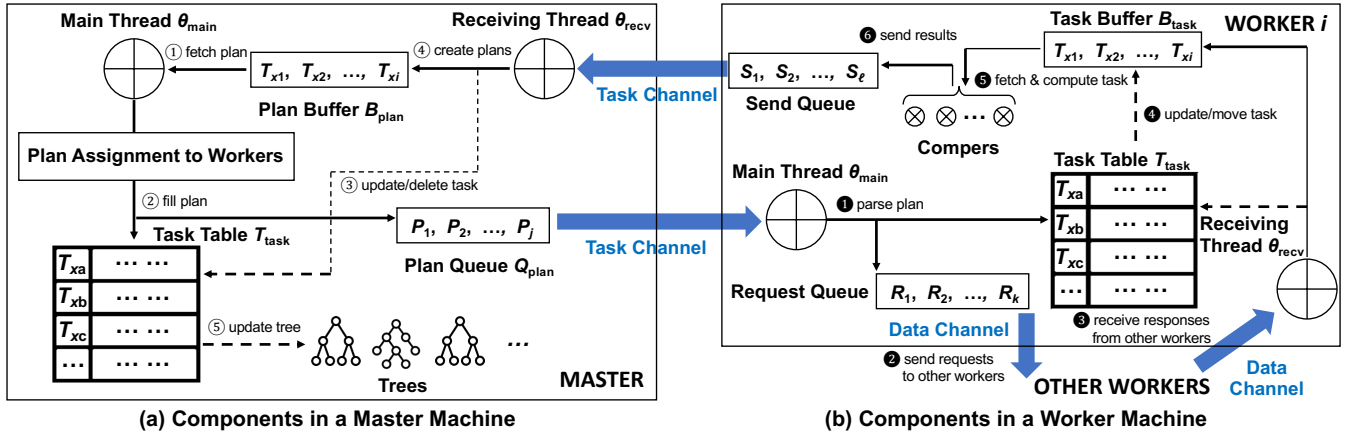


Fig. 14. Components of Master and Other Machines

that the last task of Tree  $t_{id}$  is completed so it can output the tree and update the bookkeeping information accordingly to reflect the completion of constructing Tree  $t_{id}$ .

#### D. Other TreeServer Features

TreeServer automatically infers the most appropriate algorithm (run by a column-task) to find the best split-condition of each attribute  $A_i$  based on its data type, so the input data is totally flexible. Also, while a decision tree usually maintains the predicted labels or PMF vectors only at leaf nodes, we also let internal nodes maintain them. Note that the predicted labels or PMF vectors can be easily computed as a byproduct since each node  $x$  has access to  $D_x$  during training.

When predicting the label of a new entity, this design allows the search to stop at any depth to output the predicted label at the current node. Therefore, if we train a tree with maximum depth  $d_{max}$ , we can use it for prediction as a tree with maximum depth anywhere from 1 to  $d_{max}$  without the need to train those trees.

This design also handles new attribute values unseen during training. Specifically, for a test entity, when we visit a node  $x$  whose split-condition involves attribute  $A_i$ , and find that the entity's  $A_i$ -value is never unseen in  $D_x$  during training, we then treat node  $x$  as a leaf directly to report its predicted label or PMF vector for the current entity. This is because it is unreasonable for the entity to go to either child of  $x$ .

The same applies when we encounter a missing value for attribute  $A_i$  when we visit a node  $x$  for a test entity, in which case we directly report  $x$ 's predicted label or PMF vector for the current entity.

#### E. T-thinker Components & Thread Workflow

Hereafter, we call a task  $t_x$  not yet sent to workers for processing as a **plan**. In master, all newly created plans that have not been assigned for processing are kept in the plan buffer  $B_{plan}$  (see Fig. 14(a)) implemented as a deque, while after the workers that will process a plan  $t_x$  are determined,  $t_x$  is placed into a plan queue  $Q_{plan}$  (see Fig. 14(a)) to be sent to workers for processing. Note that a new plan in  $B_{plan}$  may generate multiple column-task plans in  $Q_{plan}$  directed to different workers that collectively hold columns  $A_i \in \mathcal{C}$ .

Fig. 14 shows the components for task processing in the master (left) and worker machines (right). Besides deque  $B_{plan}$  protected with mutex, a task table  $T_{task}$  is implemented as a concurrent hash table so that the insertion and fetching of different tasks may proceed concurrently as long as they are not in the same bucket, while  $Q_{plan}$ ,  $B_{task}$  and all message queues are implemented as concurrent queues [25] that support simultaneous enqueue and dequeue operations.

We next explain the thread workflows in the master and workers, respectively.

**Thread Workflows in Master.** Master has two key threads: (1) main thread  $\theta_{main}$  and (2) receiving thread  $\theta_{recv}$ .

Thread  $\theta_{main}$  loops the following operations. It fetches a plan from  $B_{plan}$  to compute its worker assignment, i.e., operation ① in Fig. 14(a). Meanwhile, if the tasks in  $B_{plan}$  belong to less than  $n_{pool}$  trees,  $\theta_{main}$  replenishes the next unprocessed “root”-node task into to  $B_{plan}$  to start its processing.

If  $\theta_{main}$  obtains a plan  $t_x$  from  $B_{plan}$ , it computes the worker assignment using the algorithm described in Section VI, and then appends the resulting plan(s) to a plan queue  $Q_{plan}$  (see Fig. 14(a)) that sends plans to workers in batches. While a plan  $t_x$  is being processed by workers, the task information of  $t_x$  is also inserted into the master's task table  $T_{task}$  waiting for the computation results to be sent back. These correspond to operation ② in Fig. 14(a).

If  $\theta_{main}$  cannot find a plan in  $B_{plan}$ , it sleeps for  $100 \mu s$  to avoid busy waiting before probing  $B_{plan}$  again. The loop terminates when all trees have been processed and  $B_{plan}$  is empty, after which  $\theta_{main}$  flags all other threads at the master side to terminate their queue probing loop; it also inserts special messages in  $Q_{plan}$  to notify workers to terminate.

The other key thread  $\theta_{recv}$  processes each received message that contains the computed result of a task  $t_x$ , and uses it to update the table entry of  $t_x$  in  $T_{task}$ ; if  $t_x$  is determined to be finished,  $\theta_{recv}$  further updates the tree under construction using the task result, removes  $t_x$  from  $T_{task}$  and deletes the task object  $t_x$  (see operations ③ and ⑤ in Fig. 14(a)).

A completed column-task  $t_x$  will generate two child task-plans  $t_{x_\ell}$  and  $t_{x_r}$  (unless  $x$  is a leaf), and they are added to the plan buffer  $B_{plan}$  to be scheduled for processing. This process

is illustrated by operations ④ in Fig. 14(a). To prevent premature exit of  $\theta_{main}$ , we require  $\theta_{recv}$  to always add plans  $t_{x_\ell}$  and  $t_{x_r}$  into  $B_{plan}$  before decrement the tree-counter of  $t_x$  to indicate the completion of  $t_x$ , so that  $B_{plan}$  is never empty before  $t_x$ 's tree is fully constructed.

**Thread Workflows in a Worker.** As Fig. 14(b) shows, a worker runs a main thread  $\theta_{main}$  to receive task-plans from the master, a receiving thread  $\theta_{recv}$  that prepares data for tasks, and a pool of computing threads (**compers**) to compute tasks.

A task-plan is processed in 5 steps. ①:  $\theta_{main}$  keeps receiving plan-messages from the master, until a termination notification is received to terminate the worker program. For each plan-message,  $\theta_{main}$  parses it to decide if it is a subtree-task or a column-task. A subtree-task  $t_x$  still needs to obtain its data  $D_x$  from other workers, so  $\theta_{main}$  poses those data requests to a request queue for sending, as shown by operation ② in Fig. 14(b);  $\theta_{main}$  also puts  $t_x$  in the task table  $T_{task}$  waiting for the data responses. ③: when  $\theta_{recv}$  receives the data requested by  $t_x$ , it attaches the data to  $t_x$ 's entry in  $T_{task}$ ; if all necessary data are now with  $t_x$ ,  $\theta_{recv}$  will further move  $t_x$  from  $T_{task}$  to  $B_{task}$  to be fetched by a compers for processing. This process is illustrated by operation ④ in Fig. 14(b). ⑤: compers concurrently fetch tasks from  $B_{task}$  for computation; each compers keeps fetching and computing tasks until  $\theta_{main}$  flags termination, and if there is no task in  $B_{task}$ , the compers will sleep for 100  $\mu s$  to avoid busy waiting. ⑥: the computed results are appended by the compers to a sending queue which delivers them to the master in batches.

**Fault Tolerance.** Since a TreeServer program is master-driven, the master is the only single point of failure which can be strengthened by enabling a secondary master. Specifically, a worker can act as a secondary master that periodically communicates with the master to check if it is reachable. If not, the secondary master will send a special message to the other workers notifying them that it is now the new master so that they will direct task-channel messages to the secondary master. When the secondary master is enabled, the master needs to periodically synchronize the job metadata and tree construction progress to the secondary master. New tasks assigned since the last synchronization will be reassigned by the secondary master, which accepts but ignores old responses.

If a worker crashes, the master simply reassigns its lost columns to other machines by copying from the column replicas, and for each task in task table  $T_{task}$  whose computation involves the crashed worker, the master notifies workers to revoke these tasks and to delete their task objects; we also move these tasks in the master from  $T_{task}$  back to the head of  $B_{plan}$  so that their new workers can be reassigned ASAP.

#### F. Extra-Tree Support in TreeServer Deep Forests

Besides random forests, deep forest also uses “completely random decision trees” (aka. extra-trees) as the other forest type: such a tree resamples a column  $A_i$  from all attributes after each node splitting, and randomly samples a splitting value  $v$  from  $[min, max]$  where  $min$  and  $max$  are the smallest

and largest  $A_i$ -values in  $D_x$ , respectively. TreeServer also supports extra-trees, where the only difference from building random forests is that a subtree-task needs to get data from all columns (of rows  $I_x$ ) for subsequent column-sampling after each node-splitting when building the subtree, rather than simply obtain those columns of  $\mathcal{C}$  as in a random forest.

#### G. The Loan Dataset

The dataset *loan* [7] has two tables: “Origination Data” describing information of each loan, and “Monthly Performance Data” describing the monthly loan payment information. The two tables were joined on the attribute “LOAN SEQUENCE NUMBER” to obtain the final data table  $D$ . Since  $D$  has a lot of missing data, we removed every column with more than 75% missing values, and cleansed the rest by filling missing values with the mean attribute value. We also took the loan data from the latest month, the latest year, and the latest 2 years to obtain 3 datasets of different sizes as shown in Table I.